

```

# Maritime Vessel Classification - EXTREME PERFORMANCE IMPLEMENTATION
# Target: >95% Accuracy with Advanced ML Techniques
# World-Class Professional Implementation

import os
import sys
import time
import warnings
import gc
import zipfile
from pathlib import Path
from typing import Dict, List, Tuple, Any, Optional
import psutil

# Core data science libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.ensemble import RandomForestClassifier, ExtraTreesClassifier, GradientBoostingClassifier
from sklearn.model_selection import train_test_split, StratifiedKFold, cross_val_score, GridSearchCV
from sklearn.preprocessing import StandardScaler, LabelEncoder, PolynomialFeatures
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score, f1_score
from sklearn.utils.class_weight import compute_class_weight
from sklearn.feature_selection import SelectKBest, f_classif, RFE
from sklearn.linear_model import LogisticRegression
from lightgbm import LGBMClassifier
from imblearn.over_sampling import SMOTE, ADASYN
from imblearn.combine import SMOTETomek
import joblib

# Advanced ML libraries
try:
    from xgboost import XGBClassifier
    from catboost import CatBoostClassifier
    ADVANCED_MODELS_AVAILABLE = True
except ImportError:
    print("📦 Installing advanced ML libraries...")
    import subprocess
    subprocess.check_call([sys.executable, "-m", "pip", "install", "xgboost", "catboost", "-q"])
    from xgboost import XGBClassifier
    from catboost import CatBoostClassifier
    ADVANCED_MODELS_AVAILABLE = True

# Suppress warnings for cleaner output
warnings.filterwarnings('ignore')
os.environ['TF_CPP_MIN_LOG_LEVEL'] = '3'

print("🚢 MARITIME VESSEL CLASSIFICATION - EXTREME PERFORMANCE")
print("=" * 70)
print("🎯 Target: >95% Accuracy | 🏆 World-Class Implementation")
print("🔧 Advanced Feature Engineering + Hyperparameter Optimization")
print("=" * 70)

class ResourceManager:
    """Advanced resource management"""

    @staticmethod
    def get_memory_usage():
        process = psutil.Process(os.getpid())
        return process.memory_info().rss / 1024 / 1024 # MB

    @staticmethod
    def optimize_memory():
        gc.collect()
        return ResourceManager.get_memory_usage()

    @staticmethod
    def check_resources():
        memory = psutil.virtual_memory()
        print(f"📊 System Memory: {memory.total / (1024**3):.1f}GB Total, {memory.available / (1024**3):.1f}GB Available")
        print(f"📦 Current Usage: {ResourceManager.get_memory_usage():.1f}MB")
        return memory.available / (1024**3) > 1

class AdvancedDataProcessor:
    """Advanced AIS data processing with sophisticated feature engineering"""

```

```

def __init__(self, config: Dict[str, Any]):
    self.config = config
    self.vessel_type_mapping = {
        30: 'FISHING',
        31: 'TUG',
        37: 'PLEASURE',
        52: 'TUG',
        70: 'CARGO'
    }

def load_and_validate_data(self, zip_path: str) -> pd.DataFrame:
    """Load and validate AIS data"""
    print("📦 Loading AIS data...")

    with zipfile.ZipFile(zip_path, 'r') as zip_ref:
        file_list = zip_ref.namelist()
        csv_files = [f for f in file_list if f.endswith('.csv')]

        if not csv_files:
            raise ValueError("No CSV files found in ZIP archive")

        csv_file = max(csv_files, key=lambda x: zip_ref.getinfo(x).file_size)
        print(f"📄 Extracting: {csv_file}")

        with zip_ref.open(csv_file) as f:
            chunk_size = 50000
            chunks = []

            for chunk in pd.read_csv(f, chunksize=chunk_size, low_memory=False):
                chunks.append(chunk)
                if len(chunks) * chunk_size > 500000:
                    print(f"⚠️ Limited to {len(chunks) * chunk_size} records for memory efficiency")
                    break

            df = pd.concat(chunks, ignore_index=True)

        print(f"✅ Loaded {len(df):,} records")
        ResourceManager.optimize_memory()

    return self._validate_required_columns(df)

def _validate_required_columns(self, df: pd.DataFrame) -> pd.DataFrame:
    """Validate and ensure required columns exist"""
    required_cols = ['LAT', 'LON', 'SOG', 'COG', 'VesselType', 'MMSI', 'BaseDateTime']

    col_mapping = {}
    for col in required_cols:
        if col in df.columns:
            continue

        variations = {
            'LAT': ['Latitude', 'lat', 'latitude'],
            'LON': ['Longitude', 'lon', 'longitude'],
            'SOG': ['Speed', 'sog', 'speed'],
            'COG': ['Course', 'cog', 'course'],
            'VesselType': ['VesselAndCargoType', 'vessel_type'],
            'MMSI': ['mmsi'],
            'BaseDateTime': ['Timestamp', 'DateTime', 'timestamp']
        }

        for var in variations.get(col, []):
            if var in df.columns:
                col_mapping[var] = col
                break
        else:
            raise ValueError(f"Required column '{col}' not found")

    df = df.rename(columns=col_mapping)

    keep_cols = required_cols.copy()
    for col in ['Length', 'Width', 'Draft', 'length', 'width', 'draft']:
        if col in df.columns:
            keep_cols.append(col)

    return df[keep_cols].copy()

```

```

def advanced_data_cleaning(self, df: pd.DataFrame) -> pd.DataFrame:
    """Advanced data cleaning with statistical outlier detection"""
    print("🔧 Advanced data cleaning...")
    initial_rows = len(df)

    # Convert timestamp
    df['BaseDateTime'] = pd.to_datetime(df['BaseDateTime'], errors='coerce')
    df = df.dropna(subset=['BaseDateTime'])

    # Remove invalid coordinates
    df = df[
        (df['LAT'].between(-90, 90)) &
        (df['LON'].between(-180, 180)) &
        (df['LAT'] != 0) & (df['LON'] != 0)
    ]

    # Advanced speed filtering using IQR method
    Q1_speed = df['SOG'].quantile(0.25)
    Q3_speed = df['SOG'].quantile(0.75)
    IQR_speed = Q3_speed - Q1_speed
    speed_lower = max(0, Q1_speed - 1.5 * IQR_speed)
    speed_upper = min(60, Q3_speed + 1.5 * IQR_speed) # Cap at 60 knots
    df = df[df['SOG'].between(speed_lower, speed_upper)]

    # Valid course (0-359)
    df = df[df['COG'].between(0, 359)]

    # Focus on main vessel types
    df = df[df['VesselType'].isin(self.vessel_type_mapping.keys())]

    # Advanced trajectory filtering - require minimum points AND time span
    vessel_stats = df.groupby('MMSI').agg({
        'BaseDateTime': ['count', lambda x: (x.max() - x.min()).total_seconds() / 3600],
        'SOG': 'std',
        'COG': 'std'
    }).round(2)
    vessel_stats.columns = ['point_count', 'time_span_hours', 'speed_std', 'course_std']

    # Filter vessels with sufficient data quality
    valid_vessels = vessel_stats[
        (vessel_stats['point_count'] >= self.config['min_trajectory_length']) &
        (vessel_stats['time_span_hours'] >= 0.5) & # At least 30 minutes
        (vessel_stats['speed_std'] > 0.1) & # Some speed variation
        (vessel_stats['course_std'] > 1) # Some course variation
    ].index

    df = df[df['MMSI'].isin(valid_vessels)]

    print(f"✅ Advanced cleaning: {initial_rows:,} → {len(df):,} records ({100*len(df)/initial_rows:.1f}% retained)")

    df['VesselTypeClean'] = df['VesselType'].map(self.vessel_type_mapping)
    return df

def extract_world_class_features(self, df: pd.DataFrame) -> Tuple[pd.DataFrame, pd.DataFrame]:
    """Extract world-class features for >95% accuracy"""
    print("🌐 Extracting world-class features...")

    features_list = []

    for mmsi in df['MMSI'].unique():
        vessel_data = df[df['MMSI'] == mmsi].sort_values('BaseDateTime')

        if len(vessel_data) < self.config['min_trajectory_length']:
            continue

        vessel_type = vessel_data['VesselTypeClean'].iloc[0]

        # Convert to numpy arrays
        speeds = vessel_data['SOG'].values
        courses = vessel_data['COG'].values
        lats = vessel_data['LAT'].values
        lons = vessel_data['LON'].values
        timestamps = vessel_data['BaseDateTime'].values

        # Time analysis
        time_diffs = np.diff([pd.Timestamp(t).timestamp() for t in timestamps]) / 3600 # hours
        time_span = (timestamps[-1] - timestamps[0]) / np.timedelta64(1, 'h')

```

```

# === ADVANCED SPEED FEATURES ===
speed_features = self._extract_advanced_speed_features(speeds)

# === ADVANCED COURSE FEATURES ===
course_features = self._extract_advanced_course_features(courses)

# === ADVANCED GEOGRAPHIC FEATURES ===
geo_features = self._extract_advanced_geographic_features(lats, lons)

# === ADVANCED TEMPORAL FEATURES ===
temporal_features = self._extract_advanced_temporal_features(time_diffs, time_span, len(vessel_data))

# === VESSEL BEHAVIOR SIGNATURES ===
behavior_features = self._extract_behavior_signatures(speeds, courses, lats, lons, time_diffs)

# === OPERATIONAL PATTERN FEATURES ===
operational_features = self._extract_operational_patterns(speeds, courses, time_diffs)

# === STATISTICAL COMPLEXITY FEATURES ===
complexity_features = self._extract_complexity_features(speeds, courses, lats, lons)

# === PHYSICAL CHARACTERISTICS ===
physical_features = self._extract_physical_features(vessel_data)

# Combine all features
vessel_features = {
    'MMSI': mmsi,
    'VesselType': vessel_type,
    **speed_features,
    **course_features,
    **geo_features,
    **temporal_features,
    **behavior_features,
    **operational_features,
    **complexity_features,
    **physical_features
}

features_list.append(vessel_features)

features_df = pd.DataFrame(features_list)
features_df = features_df.fillna(0)

print(f"✅ Extracted {len(features_df.columns)-2} world-class features for {len(features_df)} vessels")

X = features_df.drop(['MMSI', 'VesselType'], axis=1)
y = features_df['VesselType']

return X, y

def _extract_advanced_speed_features(self, speeds):
    """Extract sophisticated speed features"""
    if len(speeds) == 0:
        return {}

    # Basic statistics
    speed_stats = {
        'speed_mean': np.mean(speeds),
        'speed_median': np.median(speeds),
        'speed_std': np.std(speeds),
        'speed_var': np.var(speeds),
        'speed_min': np.min(speeds),
        'speed_max': np.max(speeds),
        'speed_range': np.max(speeds) - np.min(speeds),
        'speed_iqr': np.percentile(speeds, 75) - np.percentile(speeds, 25),
    }

    # Advanced statistical measures
    speed_advanced = {
        'speed_skewness': self._safe_skewness(speeds),
        'speed_kurtosis': self._safe_kurtosis(speeds),
        'speed_cv': np.std(speeds) / max(np.mean(speeds), 0.1),
        'speed_gini': self._gini_coefficient(speeds),
    }

    # Percentiles and quantiles
    percentiles = [5, 10, 25, 75, 90, 95]

```

```

speed_percentiles = {f'speed_p{p}': np.percentile(speeds, p) for p in percentiles}

# Speed pattern analysis
speed_patterns = {
    'stopped_ratio': np.mean(speeds < 0.5),
    'very_slow_ratio': np.mean(speeds < 2),
    'slow_ratio': np.mean((speeds >= 2) & (speeds < 5)),
    'medium_ratio': np.mean((speeds >= 5) & (speeds < 12)),
    'fast_ratio': np.mean((speeds >= 12) & (speeds < 20)),
    'very_fast_ratio': np.mean(speeds >= 20),
    'cruise_speed_ratio': np.mean((speeds >= 8) & (speeds <= 15)),
}

# Speed change analysis
if len(speeds) > 1:
    speed_changes = np.abs(np.diff(speeds))
    speed_change_features = {
        'speed_change_mean': np.mean(speed_changes),
        'speed_change_std': np.std(speed_changes),
        'speed_change_max': np.max(speed_changes),
        'acceleration_events': np.mean(speed_changes > 2),
        'speed_stability': 1 - (np.std(speed_changes) / max(np.mean(speeds), 0.1)),
    }
else:
    speed_change_features = {
        'speed_change_mean': 0, 'speed_change_std': 0, 'speed_change_max': 0,
        'acceleration_events': 0, 'speed_stability': 1
    }

return (**speed_stats, **speed_advanced, **speed_percentiles, **speed_patterns, **speed_change_features)

def _extract_advanced_course_features(self, courses):
    """Extract sophisticated course features"""
    if len(courses) <= 1:
        return {f'course_{k}': 0 for k in ['change_mean', 'change_std', 'change_max', 'stability',
                                           'sharp_turns', 'direction_changes', 'zigzag_pattern',
                                           'circular_variance', 'turning_radius_est']}

    # Calculate course changes handling 360-degree wraparound
    course_changes = []
    for i in range(1, len(courses)):
        diff = courses[i] - courses[i-1]
        if diff > 180:
            diff -= 360
        elif diff < -180:
            diff += 360
        course_changes.append(abs(diff))

    course_changes = np.array(course_changes)

    # Basic course change statistics
    course_stats = {
        'course_change_mean': np.mean(course_changes),
        'course_change_std': np.std(course_changes),
        'course_change_max': np.max(course_changes),
        'course_change_median': np.median(course_changes),
    }

    # Course stability and patterns
    course_patterns = {
        'course_stability': 1 - (np.std(course_changes) / 180),
        'small_turns_ratio': np.mean(course_changes < 5),
        'medium_turns_ratio': np.mean((course_changes >= 5) & (course_changes < 30)),
        'sharp_turns_ratio': np.mean(course_changes >= 30),
        'very_sharp_turns_ratio': np.mean(course_changes >= 90),
        'direction_changes': np.sum(course_changes > 45),
    }

    # Advanced course analysis
    course_advanced = {
        'zigzag_pattern': self._detect_zigzag_pattern(courses),
        'circular_variance': self._circular_variance(courses),
        'turning_radius_est': self._estimate_turning_radius(course_changes),
        'course_entropy': self._calculate_entropy(course_changes, bins=36), # 10-degree bins
    }

    return (**course_stats, **course_patterns, **course_advanced)

```

```

def _extract_advanced_geographic_features(self, lats, lons):
    """Extract sophisticated geographic features"""
    if len(lats) == 0:
        return {}

    # Basic geographic bounds
    geo_basic = {
        'lat_min': np.min(lats),
        'lat_max': np.max(lats),
        'lon_min': np.min(lons),
        'lon_max': np.max(lons),
        'lat_range': np.max(lats) - np.min(lats),
        'lon_range': np.max(lons) - np.min(lons),
        'lat_center': np.mean(lats),
        'lon_center': np.mean(lons),
        'lat_std': np.std(lats),
        'lon_std': np.std(lons),
    }

    # Area and shape analysis
    geo_area = {
        'bounding_box_area': (np.max(lats) - np.min(lats)) * (np.max(lons) - np.min(lons)),
        'convex_hull_ratio': self._convex_hull_ratio(lats, lons),
        'area_efficiency': len(set(zip(np.round(lats, 3), np.round(lons, 3)))) / len(lats),
    }

    # Distance and trajectory analysis
    if len(lats) > 1:
        distances = self._calculate_distances(lats, lons)
        geo_trajectory = {
            'total_distance': np.sum(distances),
            'avg_distance_per_segment': np.mean(distances),
            'max_distance_segment': np.max(distances),
            'distance_std': np.std(distances),
            'straight_line_distance': self._haversine_distance(lats[0], lons[0], lats[-1], lons[-1]),
            'path_efficiency': self._haversine_distance(lats[0], lons[0], lats[-1], lons[-1]) / max(np.sum(distances), 0.001),
        }
    else:
        geo_trajectory = {
            'total_distance': 0, 'avg_distance_per_segment': 0, 'max_distance_segment': 0,
            'distance_std': 0, 'straight_line_distance': 0, 'path_efficiency': 1
        }

    return {**geo_basic, **geo_area, **geo_trajectory}

def _extract_advanced_temporal_features(self, time_diffs, time_span, num_points):
    """Extract sophisticated temporal features"""
    temporal_basic = {
        'time_span_hours': time_span,
        'num_positions': num_points,
        'positions_per_hour': num_points / max(time_span, 0.1),
    }

    if len(time_diffs) > 0:
        temporal_advanced = {
            'avg_time_between_positions': np.mean(time_diffs),
            'time_between_positions_std': np.std(time_diffs),
            'max_time_gap': np.max(time_diffs),
            'min_time_gap': np.min(time_diffs),
            'time_gap_cv': np.std(time_diffs) / max(np.mean(time_diffs), 0.001),
            'time_consistency': 1 - (np.std(time_diffs) / max(np.mean(time_diffs), 0.001)),
            'long_gaps_ratio': np.mean(time_diffs > np.median(time_diffs) * 2),
        }
    else:
        temporal_advanced = {
            'avg_time_between_positions': 0, 'time_between_positions_std': 0, 'max_time_gap': 0,
            'min_time_gap': 0, 'time_gap_cv': 0, 'time_consistency': 1, 'long_gaps_ratio': 0
        }

    return {**temporal_basic, **temporal_advanced}

def _extract_behavior_signatures(self, speeds, courses, lats, lons, time_diffs):
    """Extract vessel-specific behavior signatures"""
    # Fishing behavior signatures
    fishing_sig = {
        'fishing_speed_sig': np.mean(speeds < 3) * (1 + np.std(speeds)),
        'fishing_pattern_sig': self._detect_fishing_pattern(speeds, courses),
    }

```

```

        'fishing_area_revisit': self._area_revisit_pattern(lats, lons),
    }

# Cargo behavior signatures
cargo_sig = {
    'cargo_speed_sig': np.mean((speeds > 8) & (speeds < 25)) * (1 - np.std(speeds)/max(np.mean(speeds), 1)),
    'cargo_linearity_sig': self._linearity_signature(lats, lons),
    'cargo_efficiency_sig': self._movement_efficiency_signature(speeds, time_diffs),
}

# Pleasure craft signatures
pleasure_sig = {
    'pleasure_speed_sig': np.mean((speeds > 10) & (speeds < 30)),
    'pleasure_variability_sig': np.std(speeds) / max(np.mean(speeds), 1),
    'pleasure_weekend_pattern': 0, # Would need day-of-week analysis
}

# Tug boat signatures
tug_sig = {
    'tug_maneuver_sig': self._maneuverability_signature(courses),
    'tug_speed_sig': np.mean(speeds < 12) * (1 + len([i for i in range(1, len(speeds)) if abs(speeds[i] - speeds[i-1]) > 2]) / max(len(speeds), 1)),
    'tug_work_pattern_sig': self._work_pattern_signature(speeds, courses),
}

return {**fishing_sig, **cargo_sig, **pleasure_sig, **tug_sig}

def _extract_operational_patterns(self, speeds, courses, time_diffs):
    """Extract operational pattern features"""
    if len(speeds) == 0:
        return {}

    # Operating states analysis
    stopped_threshold = 0.5
    slow_threshold = 3
    working_threshold = 5

    operational = {
        'time_stopped_ratio': np.mean(speeds < stopped_threshold),
        'time_slow_ratio': np.mean((speeds >= stopped_threshold) & (speeds < slow_threshold)),
        'time_working_ratio': np.mean((speeds >= working_threshold)),
        'operational_efficiency': np.mean(speeds > 1) / max(len(speeds), 1),
    }

    # State transition analysis
    if len(speeds) > 1:
        state_changes = len([i for i in range(1, len(speeds))
                             if (speeds[i] > working_threshold) != (speeds[i-1] > working_threshold)])
        operational['state_transitions'] = state_changes / max(len(speeds), 1)
    else:
        operational['state_transitions'] = 0

    return operational

def _extract_complexity_features(self, speeds, courses, lats, lons):
    """Extract statistical complexity features"""
    complexity = {}

    # Fractal dimension approximation
    complexity['trajectory_complexity'] = self._estimate_fractal_dimension(lats, lons)

    # Entropy measures
    if len(speeds) > 0:
        complexity['speed_entropy'] = self._calculate_entropy(speeds, bins=20)
    if len(courses) > 1:
        complexity['course_entropy'] = self._calculate_entropy(courses, bins=36)

    # Autocorrelation
    if len(speeds) > 10:
        complexity['speed_autocorr'] = self._autocorrelation(speeds, lag=1)
    if len(courses) > 10:
        complexity['course_autocorr'] = self._autocorrelation(courses, lag=1)

    return complexity

def _extract_physical_features(self, vessel_data):
    """Extract physical vessel characteristics"""
    physical = {}

```

```

if 'Length' in vessel_data.columns and pd.notna(vessel_data['Length'].iloc[0]):
    length = vessel_data['Length'].iloc[0]
    physical['length'] = length
    physical['size_class'] = 1 if length < 50 else (2 if length < 150 else 3)

if 'Width' in vessel_data.columns and pd.notna(vessel_data['Width'].iloc[0]):
    width = vessel_data['Width'].iloc[0]
    physical['width'] = width
    physical['length_width_ratio'] = length / max(width, 1)
    physical['vessel_size_index'] = length * width

return physical

# === HELPER METHODS FOR ADVANCED CALCULATIONS ===

def _safe_skewness(self, data):
    """Safe skewness calculation"""
    if len(data) < 3:
        return 0
    mean = np.mean(data)
    std = np.std(data)
    if std == 0:
        return 0
    return np.mean(((data - mean) / std) ** 3)

def _safe_kurtosis(self, data):
    """Safe kurtosis calculation"""
    if len(data) < 4:
        return 0
    mean = np.mean(data)
    std = np.std(data)
    if std == 0:
        return 0
    return np.mean(((data - mean) / std) ** 4) - 3

def _gini_coefficient(self, data):
    """Calculate Gini coefficient"""
    if len(data) == 0:
        return 0
    sorted_data = np.sort(data)
    n = len(data)
    cumsum = np.cumsum(sorted_data)
    return (2 * np.sum((np.arange(1, n+1) * sorted_data))) / (n * cumsum[-1]) - (n + 1) / n

def _detect_zigzag_pattern(self, courses):
    """Detect zigzag navigation pattern"""
    if len(courses) < 4:
        return 0

    direction_changes = 0
    for i in range(2, len(courses)):
        diff1 = courses[i-1] - courses[i-2]
        diff2 = courses[i] - courses[i-1]

        # Normalize differences
        diff1 = diff1 % 360
        diff2 = diff2 % 360
        if diff1 > 180:
            diff1 -= 360
        if diff2 > 180:
            diff2 -= 360

        if diff1 * diff2 < 0 and abs(diff1) > 10 and abs(diff2) > 10:
            direction_changes += 1

    return direction_changes / max(len(courses) - 2, 1)

def _circular_variance(self, angles):
    """Calculate circular variance for course data"""
    if len(angles) == 0:
        return 0

    # Convert to radians
    angles_rad = np.radians(angles)

    # Calculate mean direction
    mean_cos = np.mean(np.cos(angles_rad))

```



```

mean_sin = np.mean(np.sin(angles_rad))

# Calculate circular variance
r = np.sqrt(mean_cos**2 + mean_sin**2)
return 1 - r

def _estimate_turning_radius(self, course_changes):
    """Estimate typical turning radius from course changes"""
    if len(course_changes) == 0:
        return 0

    significant_turns = course_changes[course_changes > 5]
    if len(significant_turns) == 0:
        return 180 # No turns, very large radius

    return 180 / max(np.mean(significant_turns), 1)

def _calculate_entropy(self, data, bins=10):
    """Calculate Shannon entropy"""
    if len(data) == 0:
        return 0

    hist, _ = np.histogram(data, bins=bins)
    hist = hist[hist > 0]
    probs = hist / np.sum(hist)
    return -np.sum(probs * np.log2(probs))

def _convex_hull_ratio(self, lats, lons):
    """Calculate ratio of convex hull area to bounding box area"""
    # Simplified convex hull approximation
    lat_range = np.max(lats) - np.min(lats)
    lon_range = np.max(lons) - np.min(lons)
    bounding_area = lat_range * lon_range

    if bounding_area == 0:
        return 1

    # Approximate convex hull area using points on boundary
    unique_points = len(set(zip(np.round(lats, 4), np.round(lons, 4))))
    approx_hull_ratio = min(1.0, unique_points / max(len(lats), 1))

    return approx_hull_ratio

def _calculate_distances(self, lats, lons):
    """Calculate distances between consecutive points"""
    distances = []
    for i in range(1, len(lats)):
        dist = self._haversine_distance(lats[i-1], lons[i-1], lats[i], lons[i])
        distances.append(dist)
    return np.array(distances)

def _haversine_distance(self, lat1, lon1, lat2, lon2):
    """Calculate Haversine distance between two points"""
    R = 6371 # Earth's radius in km

    lat1, lon1, lat2, lon2 = map(np.radians, [lat1, lon1, lat2, lon2])
    dlat = lat2 - lat1
    dlon = lon2 - lon1

    a = np.sin(dlat/2)**2 + np.cos(lat1) * np.cos(lat2) * np.sin(dlon/2)**2
    c = 2 * np.arcsin(np.sqrt(a))

    return R * c

def _detect_fishing_pattern(self, speeds, courses):
    """Detect fishing-specific movement pattern"""
    if len(speeds) < 5 or len(courses) < 5:
        return 0

    # Fishing pattern: slow speeds with irregular course changes
    slow_periods = speeds < 3
    if len(courses) > 1:
        course_changes = np.abs(np.diff(courses))
        course_changes = np.minimum(course_changes, 360 - course_changes)
        irregular_movement = course_changes > 20

    # Align arrays properly
    irregular_movement = irregular_movement[slow_periods]

```

```

    if len(irregular_movement) == len(slow_periods) - 1:
        slow_periods = slow_periods[1:] # Remove first element
    elif len(slow_periods) == len(irregular_movement) - 1:
        irregular_movement = irregular_movement[1:] # Remove first element

    if len(slow_periods) == len(irregular_movement):
        fishing_pattern = np.mean(slow_periods & irregular_movement)
    else:
        fishing_pattern = np.mean(slow_periods) * 0.5 # Fallback
else:
    fishing_pattern = np.mean(slow_periods) * 0.5

return fishing_pattern

def _area_revisit_pattern(self, lats, lons):
    """Calculate how often vessel revisits same areas"""
    if len(lats) < 10:
        return 0

    # Grid the area into cells and count revisits
    lat_bins = np.linspace(np.min(lats), np.max(lats), 10)
    lon_bins = np.linspace(np.min(lons), np.max(lons), 10)

    grid_visits = {}
    for lat, lon in zip(lats, lons):
        lat_idx = np.digitize(lat, lat_bins)
        lon_idx = np.digitize(lon, lon_bins)
        grid_key = (lat_idx, lon_idx)
        grid_visits[grid_key] = grid_visits.get(grid_key, 0) + 1

    # Calculate revisit ratio
    total_visits = sum(grid_visits.values())
    unique_cells = len(grid_visits)

    return 1 - (unique_cells / max(total_visits, 1))

def _linearity_signature(self, lats, lons):
    """Calculate how linear the trajectory is"""
    if len(lats) < 3:
        return 1

    # Calculate straight-line distance vs actual path distance
    straight_line = self._haversine_distance(lats[0], lons[0], lats[-1], lons[-1])
    actual_path = np.sum(self._calculate_distances(lats, lons))

    if actual_path == 0:
        return 1

    return straight_line / actual_path

def _movement_efficiency_signature(self, speeds, time_diffs):
    """Calculate movement efficiency signature"""
    if len(speeds) == 0 or len(time_diffs) == 0:
        return 0

    # Efficiency: consistent speed with minimal stopped time
    speed_consistency = 1 - (np.std(speeds) / max(np.mean(speeds), 0.1))
    active_time_ratio = np.mean(speeds > 1)

    return speed_consistency * active_time_ratio

def _maneuverability_signature(self, courses):
    """Calculate maneuverability signature for tugs"""
    if len(courses) < 3:
        return 0

    course_changes = np.abs(np.diff(courses))
    course_changes = np.minimum(course_changes, 360 - course_changes)

    # Tugs have moderate but frequent course changes
    moderate_turns = np.mean((course_changes > 10) & (course_changes < 90))
    frequent_changes = len(course_changes[course_changes > 5]) / max(len(course_changes), 1)

    return moderate_turns * frequent_changes

def _work_pattern_signature(self, speeds, courses):
    """Detect work pattern signature"""
    if len(speeds) < 5:

```

```

        return 0

    # Work pattern: alternating periods of movement and stopping
    speed_changes = np.abs(np.diff(speeds))
    work_cycles = len([i for i in range(1, len(speed_changes))
                       if speed_changes[i] > 1 and speed_changes[i-1] < 1])

    return work_cycles / max(len(speeds), 1)

def _estimate_fractal_dimension(self, lats, lons):
    """Estimate fractal dimension of trajectory"""
    if len(lats) < 4:
        return 1

    # Box-counting method approximation
    scales = [0.001, 0.01, 0.1] # Different box sizes
    counts = []

    for scale in scales:
        # Count boxes containing trajectory points
        lat_bins = np.arange(np.min(lats), np.max(lats) + scale, scale)
        lon_bins = np.arange(np.min(lons), np.max(lons) + scale, scale)

        occupied_boxes = set()
        for lat, lon in zip(lats, lons):
            lat_idx = np.digitize(lat, lat_bins)
            lon_idx = np.digitize(lon, lon_bins)
            occupied_boxes.add((lat_idx, lon_idx))

        counts.append(len(occupied_boxes))

    # Estimate fractal dimension from log-log slope
    if len(counts) > 1 and counts[0] != counts[-1]:
        log_scales = np.log(scales)
        log_counts = np.log(counts)
        slope = (log_counts[-1] - log_counts[0]) / (log_scales[-1] - log_scales[0])
        return max(1, min(2, -slope)) # Constrain between 1 and 2

    return 1.5 # Default moderate complexity

def _autocorrelation(self, data, lag=1):
    """Calculate autocorrelation at given lag"""
    if len(data) <= lag:
        return 0

    n = len(data)
    data = data - np.mean(data)
    autocorr = np.correlate(data[:-lag], data[lag:]) / (np.var(data) * (n - lag))
    return autocorr[0] if len(autocorr) > 0 else 0

class ExtremePerformanceEnsemble:
    """Extreme performance ensemble for >95% accuracy"""

    def __init__(self, config: Dict[str, Any]):
        self.config = config
        self.models = {}
        self.meta_model = None
        self.scaler = StandardScaler()
        self.label_encoder = LabelEncoder()
        self.feature_selector = None
        self.is_fitted = False


    def train(self, X: pd.DataFrame, y: pd.Series) -> Dict[str, Any]:
        """Train extreme performance ensemble"""
        print("🔴 Training EXTREME PERFORMANCE ensemble...")
        start_time = time.time()

        # Encode labels
        y_encoded = self.label_encoder.fit_transform(y)
        n_classes = len(self.label_encoder.classes_)

        # Feature selection - keep only the most informative features
        print("🔍 Performing feature selection...")
        selector = SelectKBest(score_func=f_classif, k=min(50, X.shape[1])) # Top 50 features
        X_selected = selector.fit_transform(X, y_encoded)
        self.feature_selector = selector

```

```


print(f" Training meta-learner...")
self.meta_model = LogisticRegression(
    max_iter=2000,
    class_weight='balanced',
    random_state=42,
    solver='lbfgs'
)
self.meta_model.fit(meta_features, y_train_balanced)

# Final evaluation
test_meta_features = self._predict_meta_features(X_test)
final_predictions = self.meta_model.predict(test_meta_features)
final_accuracy = accuracy_score(y_test, final_predictions)

# Individual model evaluations
individual_accs = {}
for name, model in self.models.items():
    pred = model.predict(X_test)
    individual_accs[name] = accuracy_score(y_test, pred)

training_time = time.time() - start_time
self.is_fitted = True


results = {
    'individual_accuracies': individual_accs,
    'meta_accuracy': final_accuracy,
    'training_time': training_time,
    'n_features_selected': X_selected.shape[1],
    'n_features_original': X.shape[1],
    'n_samples': len(X),
    'n_classes': n_classes,
    'class_names': self.label_encoder.classes_,
    'test_predictions': final_predictions,
    'test_actual': y_test,
}

print(f" EXTREME training completed in {training_time:.1f}s")
print(f" FINAL META-ENSEMBLE ACCURACY: {final_accuracy:.4f} ({100*final_accuracy:.2f}%)")

for name, acc in individual_accs.items():
    print(f"    {name}: {acc:.4f}")

return results

def _train_base_models(self, X_train, y_train, n_classes):
    """Train diverse base models with optimized hyperparameters"""

    # Random Forest with aggressive parameters
    print(f" Training optimized Random Forest...")
    self.models['rf'] = RandomForestClassifier(

```

```

self.models['rf'] = RandomForestClassifier(
    n_estimators=2000,
    max_depth=30,
    min_samples_split=2,
    min_samples_leaf=1,
    max_features='sqrt',
    class_weight='balanced_subsample',
    random_state=42,
    n_jobs=-1,
    criterion='gini'
)
self.models['rf'].fit(X_train, y_train)

# Extra Trees for diversity
print("🌲 Training Extra Trees...")
self.models['et'] = ExtraTreesClassifier(
    n_estimators=1500,
    max_depth=25,
    min_samples_split=3,
    min_samples_leaf=1,
    max_features='sqrt',
    class_weight='balanced_subsample',
    random_state=43,
    n_jobs=-1
)
self.models['et'].fit(X_train, y_train)

# LightGBM with extreme parameters
print("⚡ Training extreme LightGBM...")
self.models['lgb'] = LGBMClassifier(
    objective='multiclass',
    num_class=n_classes,
    n_estimators=3000,
    max_depth=20,
    learning_rate=0.02,
    feature_fraction=0.8,
    bagging_fraction=0.8,
    bagging_freq=5,
    min_child_samples=10,
    reg_alpha=0.1,
    reg_lambda=0.1,
    class_weight='balanced',
    random_state=42,
    verbosity=-1,
    n_jobs=-1,
    force_col_wise=True
)
self.models['lgb'].fit(X_train, y_train)

# XGBoost with optimization
print("🔧 Training extreme XGBoost...")
self.models['xgb'] = XGBClassifier(
    objective='multi:softprob',
    n_estimators=2000,
    max_depth=12,
    learning_rate=0.03,
    subsample=0.8,
    colsample_bytree=0.8,
    reg_alpha=0.1,
    reg_lambda=0.1,
    random_state=42,
    n_jobs=-1,
    verbosity=0
)
self.models['xgb'].fit(X_train, y_train)

# CatBoost for additional diversity
print("🐱 Training CatBoost...")
self.models['cat'] = CatBoostClassifier(
    iterations=1500,
    depth=10,
    learning_rate=0.05,
    class_weights=[1]*n_classes,
    random_seed=42,
    verbose=False
)
self.models['cat'].fit(X_train, y_train)

```

```

# Gradient Boosting
print("☑ Training Gradient Boosting...")
self.models['gb'] = GradientBoostingClassifier(
    n_estimators=1000,
    max_depth=8,
    learning_rate=0.1,
    subsample=0.8,
    random_state=42
)
self.models['gb'].fit(X_train, y_train)

def _generate_meta_features(self, X_train, y_train):
    """Generate meta-features using cross-validation"""
    kfold = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
    meta_features = np.zeros((len(X_train), len(self.models) * len(self.label_encoder.classes_)))

    for fold, (train_idx, val_idx) in enumerate(kfold.split(X_train, y_train)):
        X_fold_train, X_fold_val = X_train[train_idx], X_train[val_idx]
        y_fold_train = y_train[train_idx]

        feature_idx = 0
        for name, model in self.models.items():
            # Clone and train model on fold
            model_clone = type(model)(**model.get_params())
            model_clone.fit(X_fold_train, y_fold_train)

            # Get probabilities for validation set
            proba = model_clone.predict_proba(X_fold_val)

            # Store in meta-features
            n_classes = proba.shape[1]
            meta_features[val_idx, feature_idx:feature_idx+n_classes] = proba
            feature_idx += n_classes

    return meta_features

def _predict_meta_features(self, X_test):
    """Generate meta-features for test set"""
    meta_features = np.zeros((len(X_test), len(self.models) * len(self.label_encoder.classes_)))

    feature_idx = 0
    for name, model in self.models.items():
        proba = model.predict_proba(X_test)
        n_classes = proba.shape[1]
        meta_features[:, feature_idx:feature_idx+n_classes] = proba
        feature_idx += n_classes

    return meta_features

def predict(self, X: pd.DataFrame) -> np.ndarray:
    """Make predictions using the extreme ensemble"""
    if not self.is_fitted:
        raise ValueError("Model must be fitted before prediction")

    # Apply same preprocessing
    X_selected = self.feature_selector.transform(X)
    X_scaled = self.scaler.transform(X_selected)

    # Generate meta-features
    meta_features = self._predict_meta_features(X_scaled)

    # Final prediction
    predictions = self.meta_model.predict(meta_features)

    return self.label_encoder.inverse_transform(predictions)

class UltimateVisualizer:
    """Ultimate visualization for >95% results"""

    @staticmethod
    def plot_ultimate_results(results: Dict[str, Any]):
        """Create ultimate results visualization"""
        fig, axes = plt.subplots(3, 3, figsize=(24, 18))
        fig.suptitle('🚨 EXTREME PERFORMANCE MARITIME CLASSIFICATION - ULTIMATE RESULTS',
                     fontsize=18, fontweight='bold')

        # Individual model accuracies
        models = list(results['Individual Accuracies'].keys())

```

```

models = list(results['individual_accuracies'].keys())
accuracies = list(results['individual_accuracies'].values())
accuracies.append(results['meta_accuracy'])
models.append('META-ENSEMBLE')

colors = plt.cm.viridis(np.linspace(0, 1, len(models)))
bars = axes[0, 0].bar(models, accuracies, color=colors)
axes[0, 0].set_title('🔍 Model Performance Comparison', fontweight='bold', fontsize=14)
axes[0, 0].set_ylabel('Accuracy')
axes[0, 0].set_ylim(0.8, 1.0)

# Add percentage labels
for bar, acc in zip(bars, accuracies):
    height = bar.get_height()
    axes[0, 0].text(bar.get_x() + bar.get_width()/2., height + 0.005,
                    f'{acc:.1%}', ha='center', va='bottom', fontweight='bold')

axes[0, 0].tick_params(axis='x', rotation=45)
axes[0, 0].axhline(y=0.95, color='gold', linestyle='--', linewidth=2, label='95% TARGET')
axes[0, 0].legend()

# Success indicator
success_color = 'green' if results['meta_accuracy'] >= 0.95 else ('orange' if results['meta_accuracy'] >= 0.90 else 'red')
success_text = '🎯 TARGET ACHIEVED!' if results['meta_accuracy'] >= 0.95 else ('⚠️ CLOSE' if results['meta_accuracy'] >= 0.90 else ' ')

axes[0, 1].text(0.5, 0.7, f'{success_text}', fontsize=20, fontweight='bold',
               ha='center', va='center', transform=axes[0, 1].transAxes, color=success_color)
axes[0, 1].text(0.5, 0.5, f'FINAL ACCURACY', fontsize=16, fontweight='bold',
               ha='center', va='center', transform=axes[0, 1].transAxes)
axes[0, 1].text(0.5, 0.3, f'{results["meta_accuracy"]:.2%}', fontsize=24, fontweight='bold',
               ha='center', va='center', transform=axes[0, 1].transAxes, color=success_color)
axes[0, 1].axis('off')

# Feature information
axes[0, 2].text(0.1, 0.9, f'📊 FEATURE ANALYSIS', fontsize=14, fontweight='bold', transform=axes[0, 2].transAxes)
axes[0, 2].text(0.1, 0.8, f'Original Features: {results["n_features_original"]}', fontsize=12, transform=axes[0, 2].transAxes)
axes[0, 2].text(0.1, 0.7, f'Selected Features: {results["n_features_selected"]}', fontsize=12, transform=axes[0, 2].transAxes)
axes[0, 2].text(0.1, 0.6, f'Feature Reduction: {100*(1-results["n_features_selected"]/results["n_features_original"]):.1f}%',
               fontsize=12, transform=axes[0, 2].transAxes)
axes[0, 2].text(0.1, 0.5, f'Training Samples: {results["n_samples"]}', fontsize=12, transform=axes[0, 2].transAxes)
axes[0, 2].text(0.1, 0.4, f'Classes: {results["n_classes"]}', fontsize=12, transform=axes[0, 2].transAxes)
axes[0, 2].text(0.1, 0.3, f'Training Time: {results["training_time"]:.1f}s', fontsize=12, transform=axes[0, 2].transAxes)
axes[0, 2].axis('off')

# Confusion Matrix
cm = confusion_matrix(results['test_actual'], results['test_predictions'])
class_names = results['class_names']

sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=class_names, yticklabels=class_names, ax=axes[1, 0])
axes[1, 0].set_title('🔍 Confusion Matrix - Meta-Ensemble', fontweight='bold')
axes[1, 0].set_xlabel('Predicted')
axes[1, 0].set_ylabel('Actual')

# Per-class metrics
report = classification_report(results['test_actual'], results['test_predictions'],
                              target_names=class_names, output_dict=True)

metrics = ['precision', 'recall', 'f1-score']
class_metrics = {cls: [report[cls][metric] for metric in metrics] for cls in class_names}

x = np.arange(len(class_names))
width = 0.25

for i, metric in enumerate(metrics):
    values = [class_metrics[cls][i] for cls in class_names]
    axes[1, 1].bar(x + i*width, values, width, label=metric.title(), alpha=0.8)

axes[1, 1].set_title('📊 Per-Class Performance', fontweight='bold')
axes[1, 1].set_xlabel('Vessel Type')
axes[1, 1].set_ylabel('Score')
axes[1, 1].set_xticks(x + width)
axes[1, 1].set_xticklabels(class_names, rotation=45)
axes[1, 1].legend()
axes[1, 1].set_ylim(0, 1.1)

# Model accuracy distribution
all_accs = list(results['individual_accuracies'].values()) + [results['meta_accuracy']]

```

```

axes[1, 2].hist(all_accs, bins=10, alpha=0.7, color='skyblue', edgecolor='black')
axes[1, 2].axvline(results['meta_accuracy'], color='red', linestyle='--', linewidth=2, label='Meta-Ensemble')
axes[1, 2].axvline(0.95, color='gold', linestyle='--', linewidth=2, label='Target (95%)')
axes[1, 2].set_title('📊 Accuracy Distribution', fontweight='bold')
axes[1, 2].set_xlabel('Accuracy')
axes[1, 2].set_ylabel('Count')
axes[1, 2].legend()

# Performance timeline
model_names = list(results['individual accuracies'].keys()) + ['Meta']
model_accs = list(results['individual accuracies'].values()) + [results['meta_accuracy']]

axes[2, 0].plot(model_names, model_accs, 'o-', linewidth=2, markersize=8)
axes[2, 0].axhline(y=0.95, color='gold', linestyle='--', label='95% Target')
axes[2, 0].set_title('📈 Model Evolution', fontweight='bold')
axes[2, 0].set_ylabel('Accuracy')
axes[2, 0].tick_params(axis='x', rotation=45)
axes[2, 0].legend()
axes[2, 0].grid(True, alpha=0.3)

# System performance
memory_usage = ResourceManager.get_memory_usage()
axes[2, 1].text(0.1, 0.9, f"🖥️ SYSTEM PERFORMANCE", fontsize=14, fontweight='bold', transform=axes[2, 1].transAxes)
axes[2, 1].text(0.1, 0.8, f"Memory Usage: {memory_usage:.1f}MB", fontsize=12, transform=axes[2, 1].transAxes)
axes[2, 1].text(0.1, 0.7, f"Models Trained: {len(results['individual accuracies']) + 1}", fontsize=12, transform=axes[2, 1].transAxes)

performance_grade = 'A+' if results['meta_accuracy'] >= 0.95 else ('A' if results['meta_accuracy'] >= 0.90 else 'B')
axes[2, 1].text(0.1, 0.6, f"Performance Grade: {performance_grade}", fontsize=12, fontweight='bold',
               color=success_color, transform=axes[2, 1].transAxes)
axes[2, 1].text(0.1, 0.5, f"Status: Production Ready", fontsize=12, color='green', transform=axes[2, 1].transAxes)
axes[2, 1].axis('off')

# Achievement summary
achievements = []
if results['meta_accuracy'] >= 0.95:
    achievements.append("🏆 >95% Accuracy Achieved!")
if results['meta_accuracy'] >= 0.90:
    achievements.append("✅ >90% Accuracy Achieved!")
if results['training_time'] < 300:
    achievements.append("⚡ Fast Training (<5min)")
if results['n_features_selected'] < results['n_features_original']:
    achievements.append("🔍 Smart Feature Selection")

axes[2, 2].text(0.1, 0.9, f"🏆 ACHIEVEMENTS", fontsize=14, fontweight='bold', transform=axes[2, 2].transAxes)
for i, achievement in enumerate(achievements):
    axes[2, 2].text(0.1, 0.8 - i*0.1, achievement, fontsize=11, transform=axes[2, 2].transAxes)
axes[2, 2].axis('off')

plt.tight_layout()
plt.show()

# Main execution pipeline
def run_extreme_performance_pipeline():
    """Execute the extreme performance maritime vessel classification pipeline"""

    config = {
        'min_trajectory_length': 6, # Slightly reduced for more data
        'test_size': 0.15,
        'random_state': 42
    }

    print("🔍 Checking system resources...")
    ResourceManager.check_resources()

    # Mount Google Drive
    try:
        from google.colab import drive
        drive.mount('/content/drive')
        data_path = '/content/drive/MyDrive/DATA/AIS_2024_10_24.zip'
        print("☁️ Google Drive mounted successfully")
    except:
        data_path = 'data/AIS_2024_10_24.zip'
        print("💻 Running in local environment")

    if not os.path.exists(data_path):
        print(f"❌ Data file not found: {data_path}")
        return None

```



```

# Initialize extreme components
processor = AdvancedDataProcessor(config)
ensemble = ExtremePerformanceEnsemble(config)

try:
    # Advanced data processing
    df = processor.load_and_validate_data(data_path)
    df_clean = processor.advanced_data_cleaning(df)

    # World-class feature extraction
    X, y = processor.extract_world_class_features(df_clean)

    print(f"\n📊 EXTREME Dataset Summary:")
    print(f"   Features: {X.shape[1]} (world-class feature engineering)")
    print(f"   Samples: {X.shape[0]} (high-quality filtered)")
    print(f"   Classes: {y.nunique()}")
    print(f"   Class Distribution:")
    for cls, count in y.value_counts().items():
        print(f"       {cls}: {count} ({100*count/len(y):.1f}%)")

    # Train extreme ensemble
    results = ensemble.train(X, y)

    # Ultimate visualization
    UltimateVisualizer.plot_ultimate_results(results)

    # Final assessment
    ResourceManager.optimize_memory()

    print("\n" + "="*80)
    print("🏆 EXTREME PERFORMANCE PIPELINE COMPLETED!")
    print(f"🎯 FINAL META-ENSEMBLE ACCURACY: {results['meta_accuracy']:.4f} ({100*results['meta_accuracy']:.2f}%)")

    if results['meta_accuracy'] >= 0.95:
        print("🌟🌟🌟 EXCEPTIONAL SUCCESS! >95% ACCURACY ACHIEVED! 🌟🌟🌟")
        print("🏆 WORLD-CLASS MARITIME VESSEL CLASSIFICATION!")
        print("🚀 READY FOR PRODUCTION DEPLOYMENT!")
    elif results['meta_accuracy'] >= 0.90:
        print("✅ EXCELLENT PERFORMANCE! >90% Accuracy achieved!")
    else:
        print("📉 Good performance, consider data quality improvements")

    print("="*80)

    return ensemble, results

except Exception as e:
    print(f"❌ Error in extreme pipeline: {str(e)}")
    import traceback
    traceback.print_exc()
    return None, None

# Execute the extreme pipeline
if __name__ == "__main__":
    print("🚀 Starting EXTREME PERFORMANCE maritime classification...")
    model, results = run_extreme_performance_pipeline()

    if model and results:
        print("\n🏆 EXTREME PERFORMANCE MODEL READY!")
        print("💾 Save: joblib.dump(model, 'extreme_maritime_classifier.joblib')")

        if results['meta_accuracy'] >= 0.95:
            print("\n🎯🎯🎯 MISSION ACCOMPLISHED! 🎯🎯🎯")
            print("✅ >95% ACCURACY TARGET ACHIEVED!")
            print("🌟 WORLD-CLASS PERFORMANCE DELIVERED!")
        else:
            print("❌ Extreme pipeline failed. Check error messages.")

```

Installing advanced ML libraries...

MARITIME VESSEL CLASSIFICATION - EXTREME PERFORMANCE

Target: >95% Accuracy | World-Class Implementation

Advanced Feature Engineering + Hyperparameter Optimization

Starting EXTREME PERFORMANCE maritime classification...

Checking system resources...

System Memory: 12.7GB Total, 10.5GB Available

Current Usage: 1432.6MB

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

Google Drive mounted successfully

Loading AIS data...

Extracting: AIS_2024_10_24.csv

Limited to 550000 records for memory efficiency

Loaded 550,000 records

Advanced data cleaning...

Advanced cleaning: 550,000 → 93,980 records (17.1% retained)

Extracting world-class features...

Extracted 99 world-class features for 1795 vessels

EXTREME Dataset Summary:

Features: 99 (world-class feature engineering)

Samples: 1795 (high-quality filtered)

Classes: 4

Class Distribution:

TUG: 1103 (61.4%)

FISHING: 329 (18.3%)

PLEASURE: 252 (14.0%)

CARGO: 111 (6.2%)

Training EXTREME PERFORMANCE ensemble...

Performing feature selection...

Selected 50 most informative features

Applying advanced class balancing...

Training data: 3728 samples after balancing

Training optimized Random Forest...

Training Extra Trees...

Training extreme LightGBM...

Training extreme XGBoost...

Training CatBoost...

Training Gradient Boosting...

Generating meta-features...

Training meta-learner...

EXTREME training completed in 7221.3s

FINAL META-ENSEMBLE ACCURACY: 0.8815 (88.15%)

rf: 0.8296

et: 0.8259

lgb: 0.8333

xgb: 0.8333

cat: 0.8333

gb: 0.8519

